

Object Classification of Tiny Images

Mark Holt

GA Data Science

10 Dec 2014

The Dataset









- CIFAR-10

Learning Multiple Layers of Features from Tiny Images. Alex Krizhevsky, April 8 2009.

www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf


- 60000 tiny color images

- Tiny = 32 by 32 pixels
- 50000 training set
- 10000 test set
- All hand-labeled
- Contain a single dominant object
- Images of 10 objects:
 - horse, ship, truck, airplane, cat,
 - frog, deer, dog, automobile, bird



Aims

- Use machine learning techniques to build an automated image recognition system



- Learn some familiarity with Python image processing tools
- Investigate image pre-processing
- Investigate 2-D Fast Fourier Transforms
- Experience the challenges of working with images
- Acquire some experience with decision trees

Pre-processing & Feature Extraction

- Convert to grayscale
- Equalize the histogram
- Gabor Filter Banks
 - Mimic simple receptor cells in the visual cortex
 - Edge detector
 - Tunable for:
 - Orientation to the image (0, 22.5, 45, 67.5, 90, 112.5, 125 degrees)
 - Frequency (0.8 – 3.2)
 - Filter size (5 -15 pixels)
 - Resistance to image rotation, reflection
- Filter Bank Pooling
 - Resistance to image scale, and small translations

Gabor Filter Bank

Filter Size & Frequency

Band 1

Band 2

Band 3

Band 4

Band 5

Band 6

Band 7




Orientation angle to the image



Combining Filter Banks

BAND = 2 ROWS FROM THE GABOR FILTER BANKS



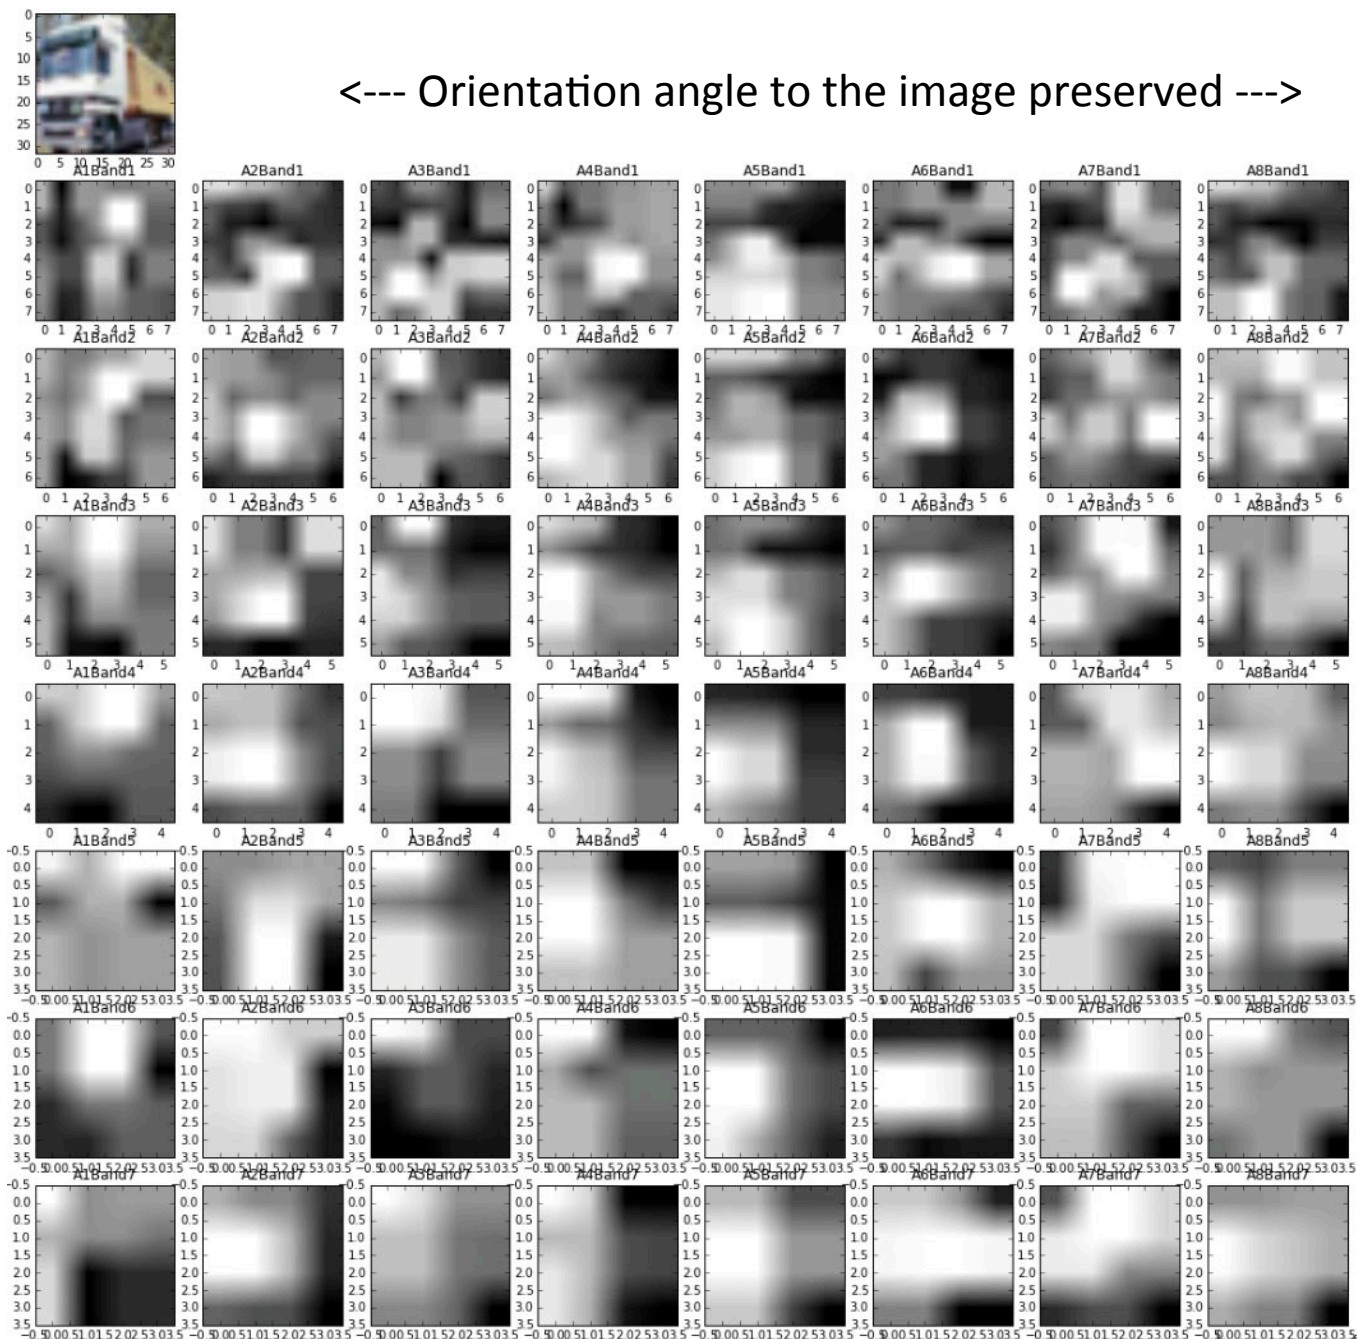
Filter Bank Pooling

<--- Orientation angle to the image preserved --->

8 by 8 pixels

Pooled Bands

- 2 closely matching layers from the Filter Banks combined



3 by 3 pixels

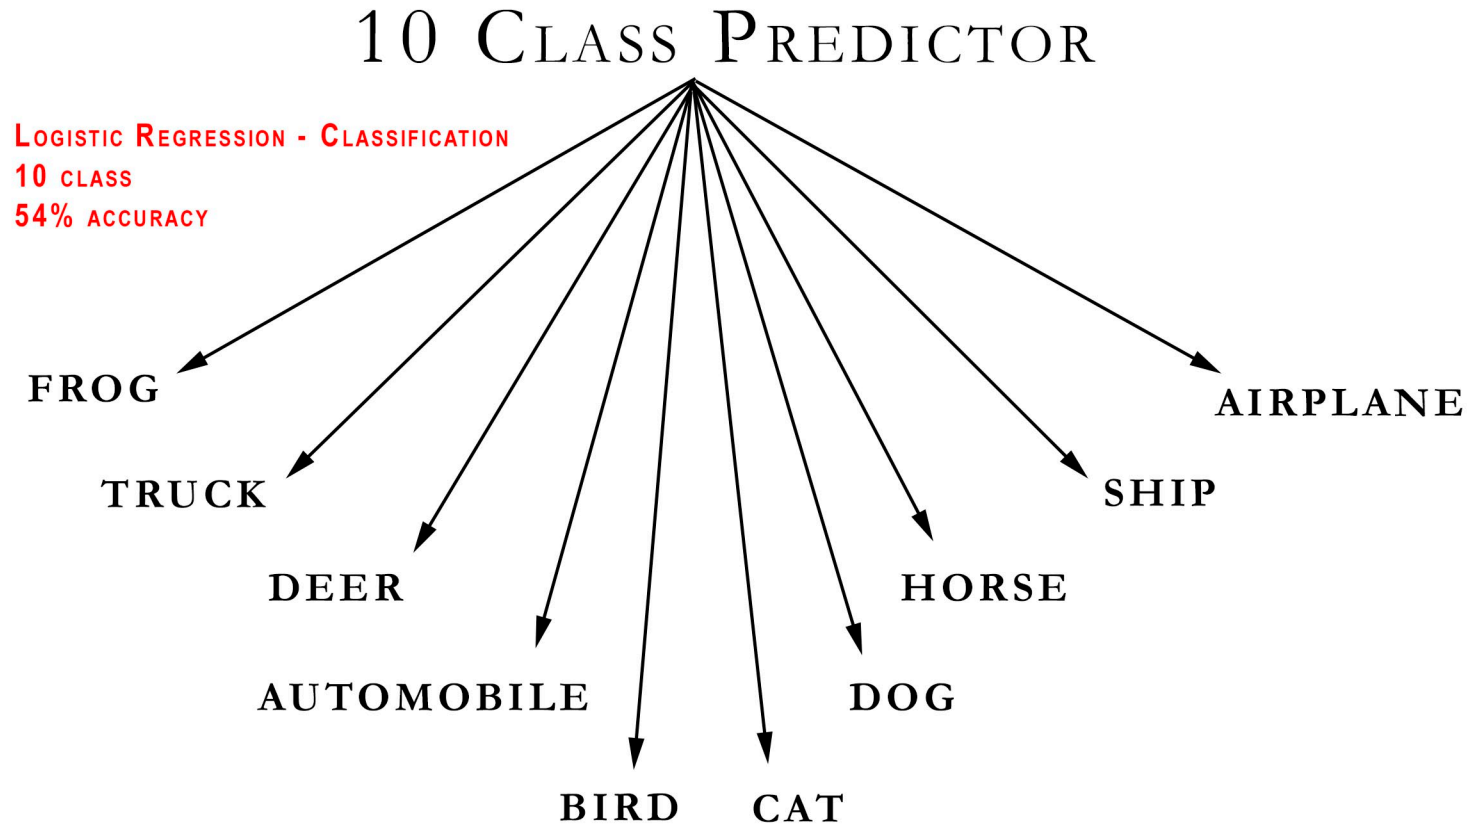
Building a Model

- Pooled Filter Banks
 - Input feature vector of length 1176
- Models investigated include:
 - Simple trees
 - Random forests
 - Extra trees
 - Stochastic gradient descent with Support Vector Machine
 - Ada Boosting
 - Logistic Regression
 - Multilayered Perceptron (MLP) – neural network
- Sheer volume of data very problematic
 - Needed to train overnight
 - Never completed
 - Kernel failed (R no better)
- Issues with python libraries
 - Opencv
 - Pybrain

Solutions

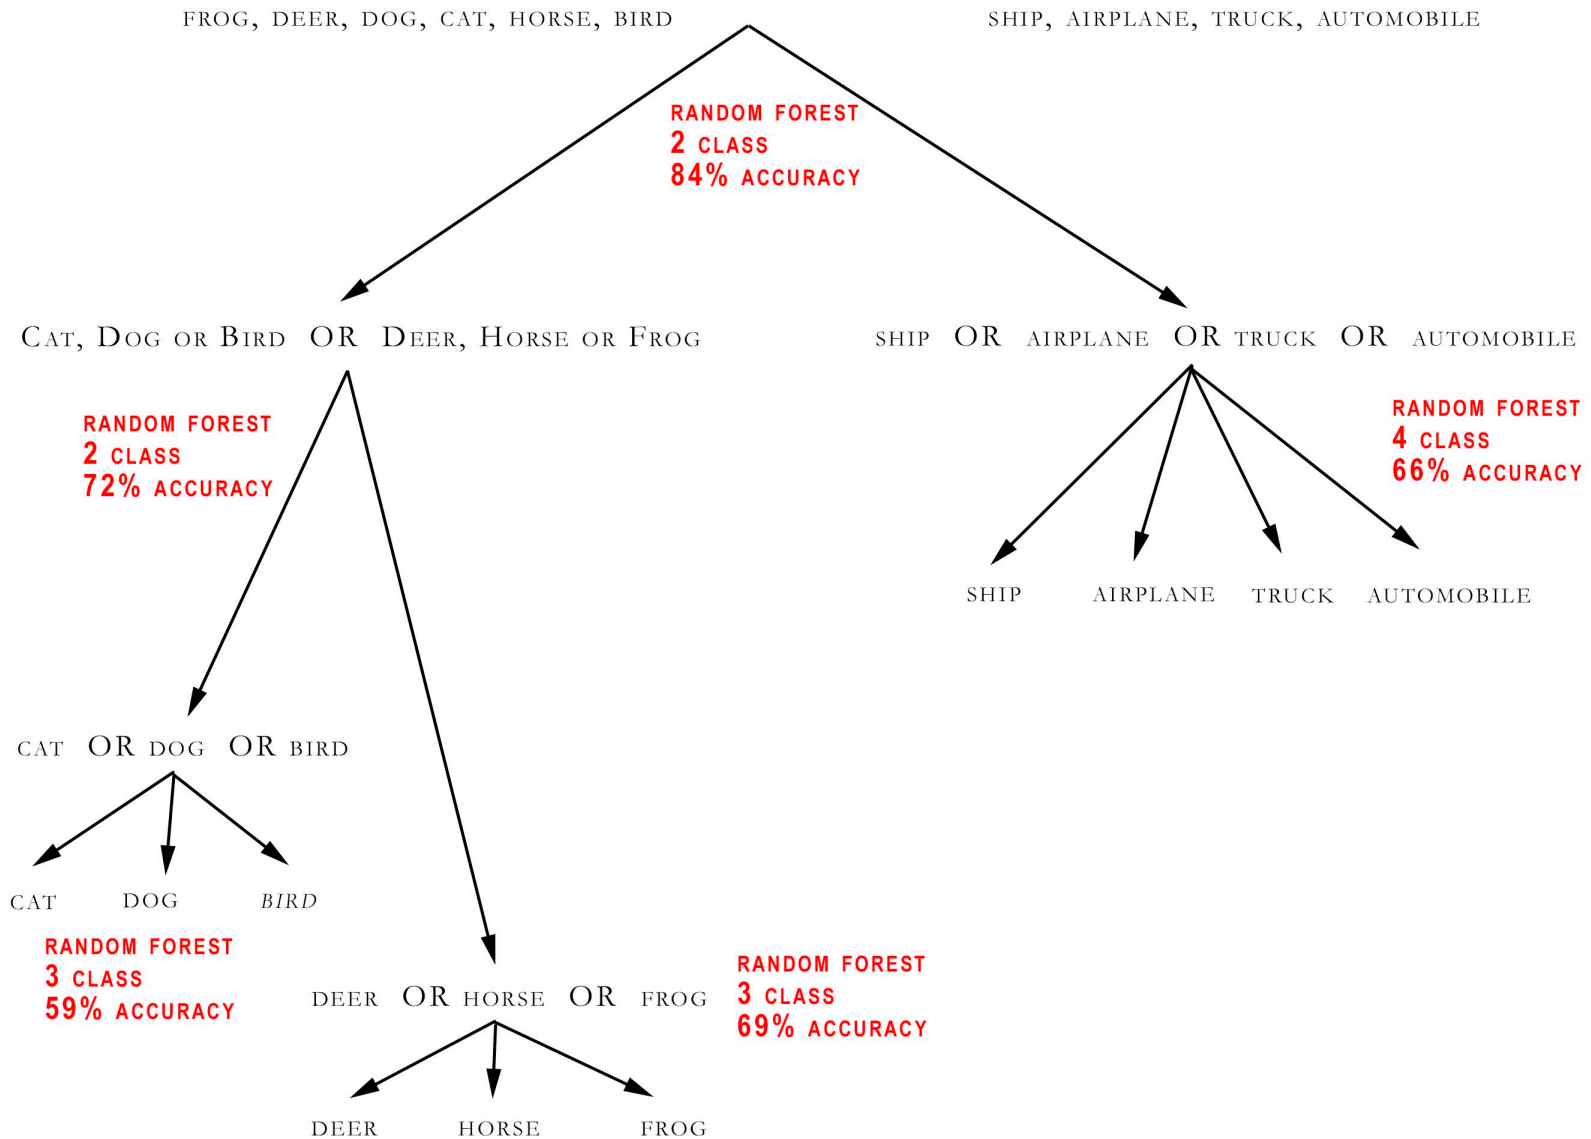
- Utilized only:
 - Logistic Regression & Random Forests
 - Random Forests, Extra Tree implementation saved the day
 - Multi-core implementation
 - Reduced processing time to minutes
 - Speed-up > 10 times
- Two Implementations:
 - Logistic Regression provided best 10 class model
 - 54% accuracy
 - Random Forests provided the best “cascade” model
 - 47% accuracy
 - State-of-the-Art
 - 82% accuracy

Classification Using a Single Model



Classification Using a Cascade

ANIMAL OR MECHANICAL



Results in More Detail

		frog	truck	deer	auto	bird	cat	dog	horse	ship	plane
Predicted	0	1	2	3	4	5	6	7	8	9	
Actual											
frog	0	608	12	51	29	81	76	67	19	18	39
truck	1	10	675	20	112	21	25	14	28	54	41
deer	2	93	22	433	35	71	86	67	127	23	43
automobile	3	39	125	19	642	18	28	14	16	61	38
bird	4	116	27	92	15	357	87	100	60	45	101
cat	5	104	36	82	38	62	349	152	72	44	61
dog	6	63	15	58	25	62	146	511	74	24	22
horse	7	21	47	75	23	47	49	79	620	16	23
ship	8	16	48	17	67	22	29	14	16	678	93
plane	9	34	46	48	43	78	36	26	25	142	522

Difficult to separate:

Deer & horse

Automobile & truck

Cats & Dogs

Cats & frogs

Birds – generally hard,

Frogs, dogs & airplanes!

“Easiest” to separate:

Ships & trucks

Demonstration

- Ten classifier in action
- Cascade classifier in action

Future Work

- Additional feature extraction
 - Deep networks
 - Prototypes
 - Auto-associative networks?
- Utilize the color channels
- Implementation for generating the models
 - Fast, low-level, GPU
- Augment the training set
 - Reflections
 - Scaling
 - Rotation
- Dimensional Reduction?
 - Simple experiments suggest little impact